



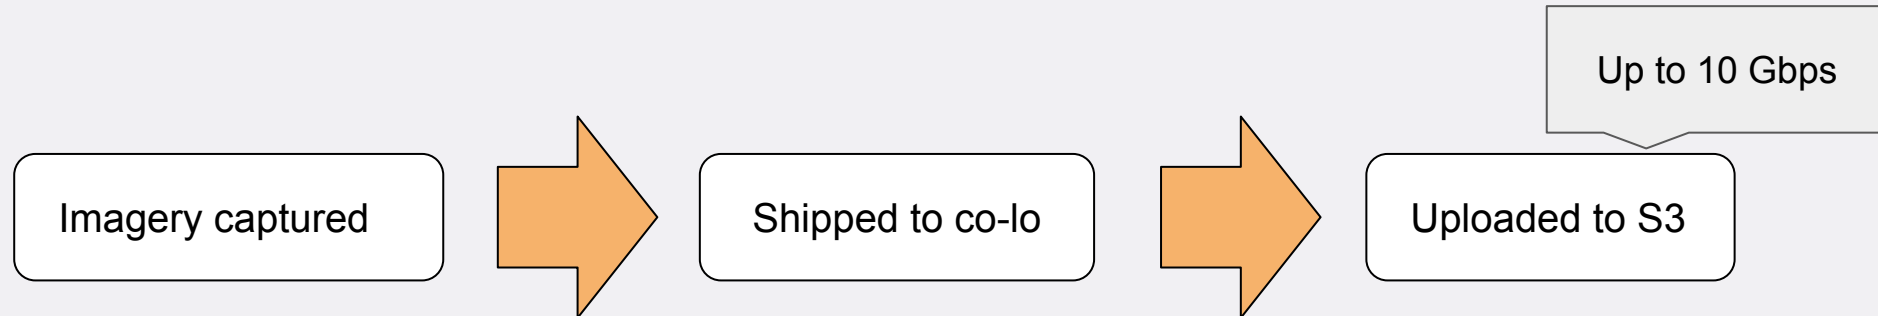
Using Terraform with AWS Batch for data processing at scale

Terraform and AWS Batch


Data processing at scale



Data ingestion



Processing



- Different resource requirements
- Different scaling characteristics



Pipeline v1

- Use existing python scripts
- Refactor to accept S3 in and out urls
- Wrap each in docker
- Terraform to setup AWS Batch
- Python script to kick off AWS Batch jobs



Terraform - set up

- From scratch in Terraform
- Just a **main.tf** and a **variables.tf**
- About 300 lines of config
- So what does that actually look like?



```
provider "aws" {  
  region = var.region  
  allowed_account_ids = var.allowed_account_ids  
}
```




```
terraform {  
  backend "s3" {  
    bucket = "bucket"  
    key = "prod/pipeline.tfstate"  
    dynamodb_table = "terraform-state-lock"  
    encrypt = true  
    region = "ap-southeast-2"  
  }  
}
```



```
resource "aws_ecr_repository" "pipeline" {  
  name = "pipeline"  
}
```



```
resource "aws_ecr_repository" "pipeline" {  
  name = "pipeline"  
}  
  
output "repository_url" {  
  description = "The URL for the pipeline ECR repository"  
  value = aws_ecr_repository.pipeline.repository_url  
}
```




```
data "aws_iam_policy_document" "pipeline_job_assume" {  
  ...  
}  
  
resource "aws_iam_role" "pipeline_job" {  
  assume_role_policy = data.aws_iam_policy_document.pipeline_job_assume.json  
}
```



```
data "aws_iam_policy_document" "pipeline_job" {  
  ...  
}  
  
resource "aws_iam_role_policy" "pipeline_job" {  
  role = aws_iam_role.pipeline_job.name  
  policy = data.aws_iam_policy_document.pipeline_job.json  
}
```



```
resource "aws_iam_role_policy_attachment" "pipeline_job" {  
  role = aws_iam_role.pipeline_job.name  
  policy_arn = ".../service-role/AmazonEC2ContainerServiceforEC2Role"  
}
```




```
resource "aws_iam_instance_profile" "pipeline_job" {  
  role = aws_iam_role.pipeline_job.name  
}
```



```
resource "aws_security_group" "pipeline" {  
  name = "pipeline-security-group"  
  
  egress {  
    from_port = 0  
    to_port   = 0  
    protocol  = "-1"  
    cidr_blocks = ["0.0.0.0/0"]  
  }  
}
```



```
data "aws_ssm_parameter" "ecs_amazon_linux_2" {  
  name = "/aws/service/ecs/optimized-ami/amazon-linux-2/recommended"  
}  
  
resource "aws_launch_template" "pipeline" {  
  name = "pipeline"  
  ebs_optimized = true  
  image_id = jsondecode(data.aws_ssm_parameter.ecs_amazon_linux_2.value).image_id  
}
```




```
resource "aws_batch_compute_environment" "pipeline" {
  compute_environment_name_prefix = "pipeline"

  compute_resources {
    instance_role = aws_iam_instance_profile.pipeline_job.arn
    instance_type = ["p3.8xlarge", "m5.16xlarge",]

    max_vcpus = 96
    min_vcpus = 0

    security_group_ids = [aws_security_group.pipeline.id,]
    subnets = [var.subnet_id,]

    launch_template {
      launch_template_id = aws_launch_template.pipeline.id
      version = aws_launch_template.pipeline.latest_version
    }

    type = "EC2"
  }

  service_role = aws_iam_role.pipeline_service.arn
  type = "MANAGED"

  lifecycle {create_before_destroy = true}
  depends_on = [aws_iam_role_policy_attachment.pipeline_service]
}
```



```
resource "aws_batch_job_queue" "stitcher" {  
  name = "stitcher"  
  state = "ENABLED"  
  priority = 1  
  compute_environments = [aws_batch_compute_environment.pipeline.arn]  
}
```



```
resource "aws_batch_job_definition" "stitcher" {  
  name = "stitcher"  
  type = "container"  
  container_properties = jsonencode({  
    image = "${aws_ecr_repository.pipeline.repository_url}:1.0"  
    command = ["stitcher"]  
    vcpus = 10  
    memory = 50000  
    resourceRequirements = [  
      {  
        type = "GPU"  
        value = "2"  
      },  
    ]  
  })  
}
```



Make it go now

```
$ terraform apply
...
Outputs:

repository_url = <url>

$ docker tag <image_id> <url>:1.0
$ docker push <url>:1.0
$ aws batch submit-job --job-name="Foo" --job-queue="stitcher" --job-definition="stitcher"
```



Lessons

- Get up and running with Terraform and Batch quickly
- AWS launch template not well supported
- External dependencies are the enemy
- Batch status monitoring and debugging is poor
- Batch limits interfaces between jobs to cli options and success or fail status.
- Batch limited to 20 dependant jobs



Thank you. Any questions?

craig@dendra.io
@craigloftus on DigitalOxford